

# Design and Analysis of Data Structures for Dynamic Trees

Renato Werneck  
Advisor: Robert Tarjan

# Outline

⇒ The Dynamic Trees problem

- Existing data structures
- A new worst-case data structure
- A new amortized data structure
- Experimental results
- Final remarks

# The Dynamic Trees Problem

- Dynamic trees:
  - Goal: maintain an  $n$ -vertex forest that changes over time.
    - \* **link**( $v, w$ ): adds edge between  $v$  and  $w$ .
    - \* **cut**( $v, w$ ): deletes edge  $(v, w)$ .
  - Application-specific data associated with vertices/edges:
    - \* updates/queries can happen in bulk (entire paths or trees at once).
- Concrete examples:
  - Find minimum-weight edge on the path between  $v$  and  $w$ .
  - Add a value to every edge on the path between  $v$  and  $w$ .
  - Find total weight of all vertices in a tree.
- Goal:  $O(\log n)$  time per operation.

# Applications


- Subroutine of **network flow algorithms**
  - maximum flow
  - minimum cost flow
- Subroutine of **dynamic graph algorithms**
  - dynamic biconnected components
  - dynamic minimum spanning trees
  - dynamic minimum cut
- Subroutine of **standard graph algorithms**
  - multiple-source shortest paths in planar graphs
  - online minimum spanning trees
- ...

# Application: Online Minimum Spanning Trees

- Problem:
  - Graph on  $n$  vertices, with new edges “arriving” one at a time.
  - Goal: maintain the minimum spanning forest (MSF) of  $G$ .
- Algorithm:
  - Edge  $e = (v, w)$  with length  $\ell(e)$  arrives:
    1. If  $v$  and  $w$  in different components: insert  $e$ ;
    2. Otherwise, find longest edge  $f$  on the path  $v \cdots w$ :
      - \*  $\ell(e) < \ell(f)$ : remove  $f$  and insert  $e$ .
      - \*  $\ell(e) \geq \ell(f)$ : discard  $e$ .


## Example: Online Minimum Spanning Trees

- Current minimum spanning forest.




## Example: Online Minimum Spanning Trees

- Edge between different components arrives.



## Example: Online Minimum Spanning Trees


- Edge between different components arrives:
  - add it to the forest.






## Example: Online Minimum Spanning Trees

- Edge within single component arrives.




## Example: Online Minimum Spanning Trees

- Edge within single component arrives:
  - determine path between its endpoints.




## Example: Online Minimum Spanning Trees

- Edge within single component arrives:
  - find longest edge on the path between its endpoints.



## Example: Online Minimum Spanning Trees

- Edge within single component arrives:
  - if longest edge on the path longer than new edge, swap them.



# Online Minimum Spanning Trees

- Data structure must support the following operations:
  - add an edge;
  - remove an edge;
  - decide whether two vertices belong to the same component;
  - find the longest edge on a specific path.
- Each in  $O(\log n)$  time:
  - basic strategy: map arbitrary tree onto balanced tree.

# Data Structures for Dynamic Trees

- Different applications require different operations.
- Desirable features of a data structure for Dynamic Trees:
  - low overhead (fast);
  - simple to implement;
  - intuitive interface (easy to adapt, specialize, modify);
  - general:
    - \* path and tree queries;
    - \* no degree constraints;
    - \* rooted and unrooted trees.

# Outline

- The Dynamic Trees problem

⇒ Existing data structures

- A new worst-case data structure
- A new amortized data structure
- Experimental results
- Final remarks


# Main Strategies

- Path decomposition:
  - ST-trees [Sleator and Tarjan 83, 85];
    - \* also known as link-cut trees or dynamic trees.
- Tree contraction:
  - Topology trees [Frederickson 85];
  - Top trees [Alstrup, Holm, de Lichtenberg, and Thorup 97];
  - RC-trees [Acar, Blelloch, Harper, Vitter, and Woo 04].
- Linearization:
  - ET-trees [Henzinger and King 95, Tarjan 97];
  - Less general: cannot handle path queries.




# Path Decomposition

- ST-trees [Sleator and Tarjan 83, 85]:
  - partition the tree into vertex-disjoint paths;
  - represent each path as a binary tree;
  - “glue” the binary trees appropriately.



# Path Decomposition

- ST-trees [Sleator and Tarjan 83, 85]:
  - partition the tree into vertex-disjoint paths;
  - represent each path as a binary tree;
  - “glue” the binary trees appropriately.
- Complexity:
  - with ordinary balanced trees:  $O(\log^2 n)$  worst-case;
  - with globally biased trees:  $O(\log n)$  worst-case;
  - with splay trees:  $O(\log n)$  amortized, but much simpler.
- Main features:
  - relatively low overhead (one node per vertex/edge);
  - adapting to different applications requires knowledge of inner workings;
  - tree-related queries require bounded degrees (or ternarization).





# Contractions: Rake and Compress

- Proposed by Miller and Reif [1985] (parallel setting).



- Rake:**

- eliminates a **degree-one** vertex;
- edge combined with successor:
  - \* assumes **circular order**.



- Compress:**

- eliminates a **degree-two** vertex;
- combines two edges into one.




- Original and resulting edges are **clusters**:

- cluster represents both a path and a subtree;
- user defines what to store in the new cluster.


# Contractions

- Contraction:
  - series of rakes and compresses;
  - reduces tree to a single cluster (edge).
- Any order of rakes and compresses is “right”:
  - final cluster will have the correct information;
  - data structure decides which moves to make:
    - \* just “asks” the user how to update values after each move.
- Example:
  - work in rounds;
  - perform a maximal set of independent moves in each round.


## Contractions: Example




## Contractions: Example



## Contractions: Example




## Contractions: Example







## Contractions: Example




## Contractions: Example




## Contractions: Example




## Contractions: Example



## Contractions: Example



## Contractions: Example








## Contractions: Example









# Top Trees: Example

en












# Top Trees: Example





# Top Trees: Example



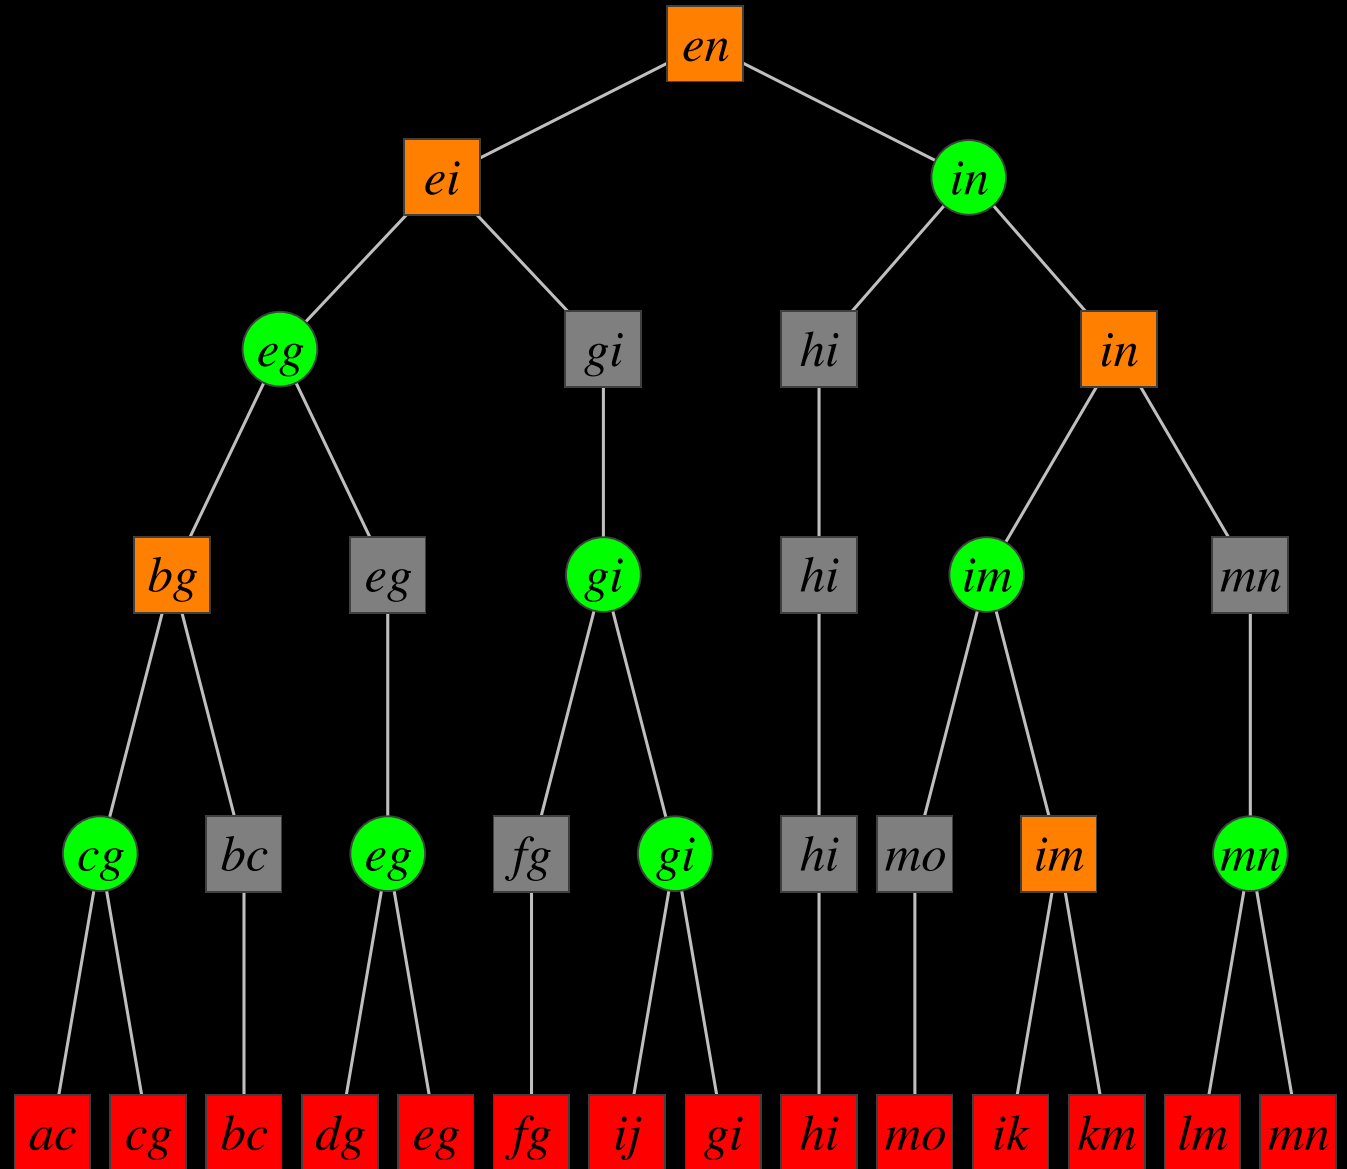

# Top Trees: Example



# Top Trees: Example





# Top Trees: Example



## Contractions: Updating Values



- To find the minimum-cost edge on the tree:

- rake:  $C \leftarrow \min\{A, B\}$
- compress:  $C \leftarrow \min\{A, B\}$



- To find the maximum edge on a path:


- rake:  $C \leftarrow B$ .
- compress:  $C \leftarrow \max\{A, B\}$



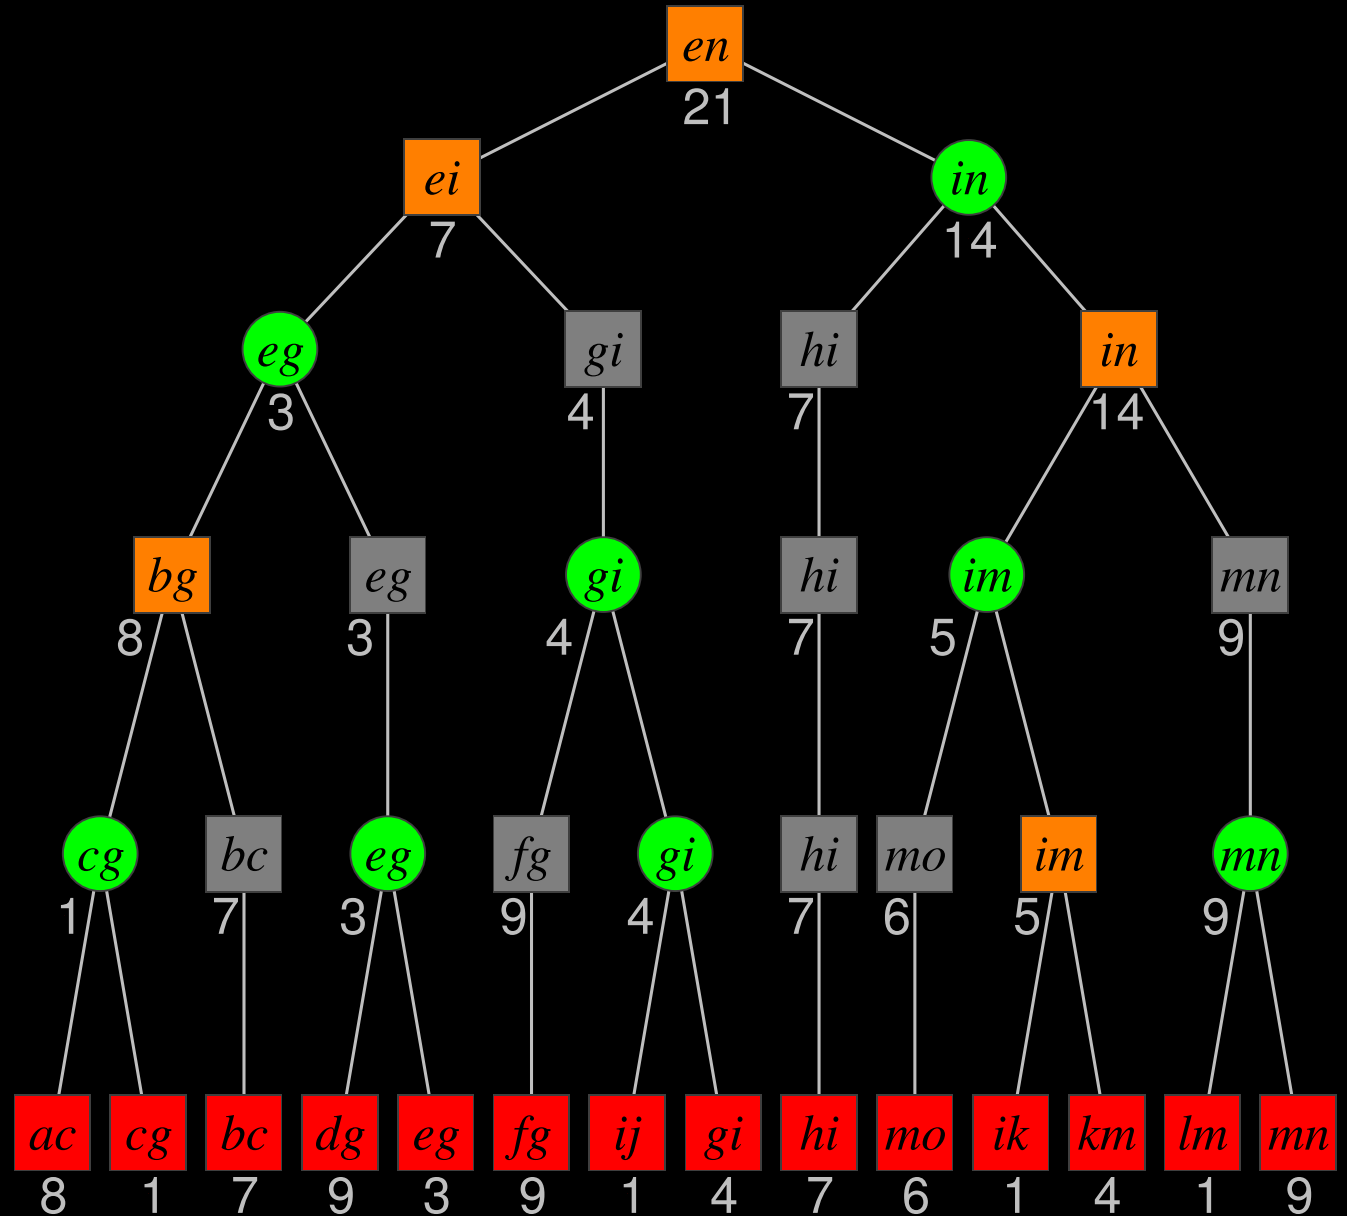

- To find the length of a path:

- rake:  $C \leftarrow B$
- compress:  $C \leftarrow A + B$

## Contractions: Example with Values



# Top Trees: Path Lengths





# Top Trees

- Top tree embodies a contraction:
  - top tree root represents the entire original tree;
  - **expose**( $v, w$ ) ensures that path  $v \cdots w$  is represented at the root;
  - interface only allows direct access to the root.
- Other operations:
  - **link**: joins two top trees;
  - **cut**: splits a top tree in two;
  - Goal: make **link**, **cut**, and **expose** run in  $O(\log n)$  time.
- Known implementation [Alstrup et al. 97]:
  - interface to **topology trees**:
    - \* variant where clusters are vertices;
    - \* degree must be bounded.

# Outline

- The Dynamic Trees problem
- Existing data structures

⇒ A new worst-case data structure

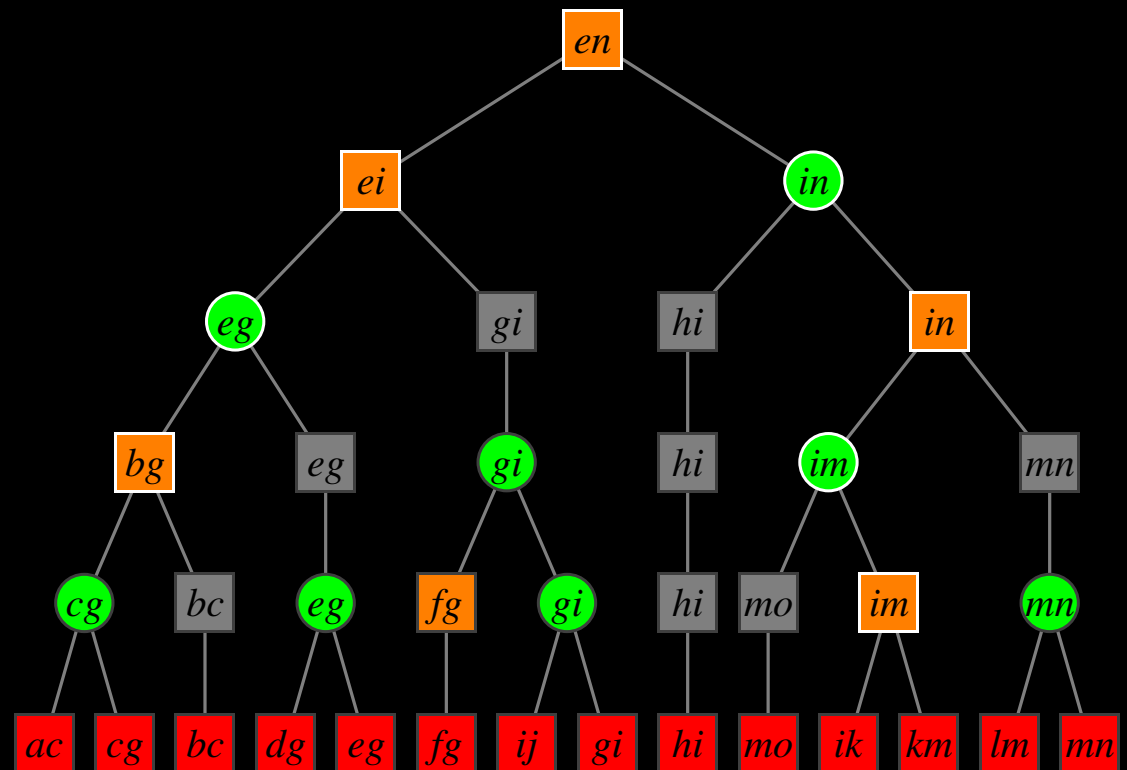
- A new amortized data structure
- Experimental results
- Final remarks

# Tree Contraction

- Contraction scheme:
  - Work in rounds, each with a maximal set of independent moves.
- Lemma: there will be at most  $O(\log n)$  levels.
  - At least **half** the vertices have degree 1 or 2.
  - At least **one third** of those will disappear:
    - \* a move blocks at most 2 others.
  - No more than  $5/6$  of the original vertices will remain.


# Online MSF on Augmented Stars

- Corollary: **expose**( $v, w$ ) can be implemented in  $O(\log n)$  time.
  - Temporarily eliminate all clusters with  $v$  or  $w$  as internal vertices:
    - \* at most two per level:  $O(\log n)$ .



# Online MSF on Augmented Stars



- Corollary: **expose**( $v, w$ ) can be implemented in  $O(\log n)$  time.
  - Temporarily eliminate all clusters with  $v$  or  $w$  as internal vertices:
    - \* at most two per level:  $O(\log n)$ .
  - Build a temporary top tree on the remaining root clusters.
  - Restore original tree.





# Tree Contraction

- Update scheme:
  - Goal: minimize “damage” to original contraction after **link** or **cut**.
  - For each level (bottom-up), execute two steps:
    1. **replicate** as many original moves as possible;
    2. perform **new moves** until maximality is achieved.
  - **Step 1 is implicit.**

## Updates: Example




## Updates: Example






## Updates: Example





these moves cannot be replicated





new move



## Updates: Example





## Updates: Example





## Updates: Example



## Updates: Example





## Updates: Example





## Updates: Example

## Updates: Example









## Updates: Example





## Updates: Example





## Updates: Example





## Updates: Example





## Updates: Example






## Updates: Example



## Updates: Example






## Updates: Example







## Updates: Example




## Updates: Example




# Running Time

- Theorem: each level can be updated in  $O(1)$  time.
- Only need to worry about the **core**:
  - connected subgraph induced by the new (**active**) clusters;
  - remaining clusters: **inactive** subgraph.



# Running Time

- Theorem: each level can be updated in  $O(1)$  time.
- Only need to worry about the **core**:
  - connected subgraph induced by the new (**active**) clusters;
  - remaining clusters: **inactive** subgraph.



# Running Time

- Theorem: each level can be updated in  $O(1)$  time.
- Only need to worry about the **core**:
  - connected subgraph induced by the new (**active**) clusters;
  - remaining clusters: **inactive** subgraph.
- Suffices to prove that **core size**  $s$  is bounded by a **constant**.
  - If core were a free tree, would shrink by  $1/6$  between rounds.
  - Each point of contact will add  $O(1)$  clusters to the core.
  - Claim: there are at most four points of contact.
  - Constant initial size + multiplicative decrease + additive increase:  
 $\Rightarrow$  constant size.

# Contraction-based Data Structure

- Positive aspects:
  - general (top tree interface);
  - conceptually simple algorithm;
  - direct implementation (does not use topology trees);
  - $O(\log n)$  worst case.
- Potential overheads:
  - pointers within and among levels:
    - \* need Euler tour of each level;
  - unmatched clusters are repeated (dummy nodes).
- Joint work with J. Holm, R. Tarjan, and M. Thorup.

# Outline


- The Dynamic Trees problem
- Existing data structures
- A new worst-case data structure

⇒ A new amortized data structure

- Experimental results
- Final remarks

# Representation


- Consider some unrooted tree:






# Representation

- Make it a **unit tree**:
  - directed tree with degree-one root.





# Representation

- Pick a **root path**:
  - starts at a leaf;
  - ends at the root.





# Representation

- Represent the root path as a **binary tree**:
  - leaves: base clusters (original edges);
  - internal nodes: compress clusters.





# Representation

- What if a vertex has degree greater than two?
  - Recursively represent each subtree rooted at the vertex (at most two, because of circular order).





# Representation

- What if a vertex has degree greater than two?
  - Recursively represent each subtree rooted at the vertex.
  - Before vertex is compressed, rake subtrees onto adjacent cluster.




# Representation

- Representation:
  - Up to four children per node (two proper ones, two foster ones)
  - Meaning: up to two rakes followed by a compress.
  - Example:  $N_y = \text{compress}(\text{rake}(D, N_x), \text{rake}(E, yz)) = wz$ .




# Representation

- How does the recursive representation work?
  - Must represent subtrees rooted at the root path.



# Representation


- How does the recursive representation work?
  - Must represent subtrees rooted at the root path.






# Representation

- How does the recursive representation work?
  - Must represent subtrees rooted at the root path.
  - Each subtree is a sequence of unit trees with a common root.





# Representation

- How does the recursive representation work?
  - Must represent subtrees rooted at the root path.
  - Each subtree is a sequence of unit trees with a common root.
  - Represent each recursively.



# Representation

- How does the recursive representation work?
  - Must represent subtrees rooted at the root path.
  - Each subtree is a sequence of unit trees with a common root.
  - Represent each recursively.
  - Build a binary tree of rakes.

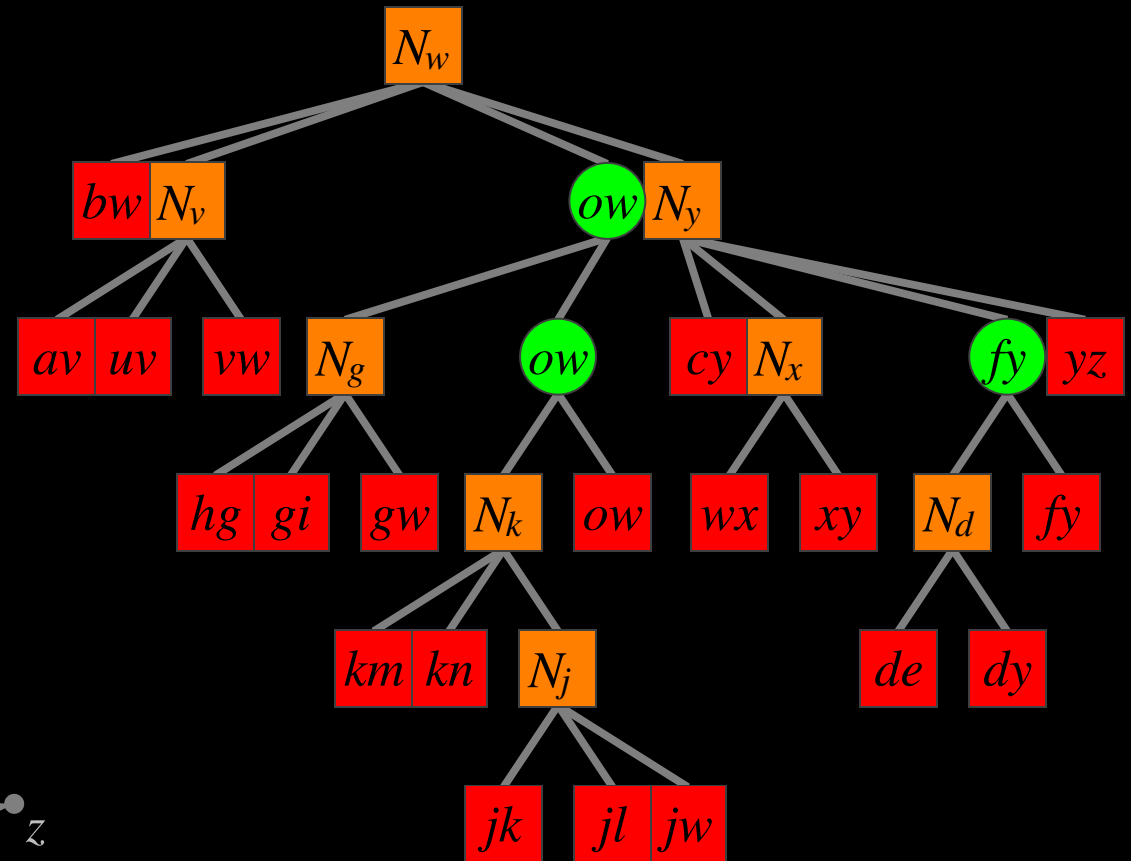



# Representation

- Two different views:
  - User interface: tree contraction.
    - \* sequence of rakes and compresses;
    - \* a single tree (a top tree).
  - Implementation: path decomposition.
    - \* maximal edge-disjoint paths;
    - \* hierarchy of binary trees (rake trees/compress trees);
    - \* similar to ST-trees.

# Representation

- Full example:



# Self-Adjusting Top Trees

- Topmost compress tree represents the root path:
  - determined by **expose**( $v, w$ );
  - implementation similar to ST-trees;
  - basic operations:
    - \* **splay**: rebalances each binary tree;
    - \* **splice**: changes the partition into paths.
- Main result: expose takes  $O(\log n)$  amortized time.
- Operations **link** and **cut** use expose as the main subroutine.
- Joint work with R. Tarjan [SODA'05].

# Outline

- The Dynamic Trees problem
- Existing data structures
- A new worst-case data structure
- A new amortized data structure

⇒ Experimental results

- Final remarks

# Experimental Results

- Data structures implemented (in C++):
  - **TOP-W**: worst-case top trees;
  - **TOP-S**: self-adjusting top trees;
  - **ET-S**: self-adjusting ET-trees;
  - **ST-V/ST-E**: self-adjusting ST-trees;
  - **LIN-V/LIN-E**: “obvious” linear-time data structure.
- Machine: 2.0 GHz Pentium IV with 1 GB of RAM.
- Applications:
  - random operations;
  - maximum flows;
  - online minimum spanning trees;
  - shortest paths.




# Experimental Results: Executive Summary

- **ST-V**: fastest  $O(\log n)$  algorithm;
  - **ST-E** and **ET-S**: up to twice as slow.
- **TOP-S** is 2–4 times as slow as **ST-V**.
- **TOP-W** vs. **TOP-S**:
  - **TOP-S** faster for links and cuts;
  - **TOP-W** faster for expose;
  - **TOP-S** benefits from correlated operations (splaying).
- **LIN-V**/**LIN-E**: fastest for paths with up to  $\sim 1000$  vertices.

# Experiment: Random Operations

- Setup:
  - Start with empty graph on  $n$  vertices;
  - use  $n - 1$  links to create a random spanning tree;
  - alternate cuts and links until  $\#ops = 10n$ .




## Experiment: Maximum Flow

- Maximum flow:
  - Given: directed graph with  $n$  vertices and  $m$  edges, source  $s$ , sink  $t$ ;
  - Goal: find maximum flow from  $s$  to  $t$ .
- Algorithm:
  - Find **shortest path** from  $s$  to  $t$  with positive residual capacity;
  - send as much flow as possible, update residual capacities, repeat;
  - running time:  $O(mn^2)$ .
- With dynamic trees:
  - Keep “partial paths” between iterations as a forest;
  - allows augmentations in  $O(\log n)$  time;
  - running time:  $O(mn \log n)$ .


# Maximum Flow on Layered Networks

- Layered networks:
  - 4 rows,  $\lfloor n/4 \rfloor$  columns between  $s$  and  $t$ ;
  - each vertex connected to 3 random vertices in the next column;
  - augmenting paths have  $\Omega(n)$  arcs.



# Maximum Flow on Square Meshes

- Square meshes:
  - full  $\lfloor \sqrt{n} \rfloor \times \lfloor \sqrt{n} \rfloor$  grid between  $s$  and  $t$ ;
  - each grid vertex connected to all 4 neighbors;
  - augmenting paths have  $\Omega(\sqrt{n})$  arcs.




# Experiment: Online Minimum Spanning Forest

- Online minimum spanning forest:
  - Edges processed one at a time;
  - Edge  $(v, w)$  inserted if
    - \*  $v$  and  $w$  in different components; or
    - \*  $(v, w)$  is shorter than longest edge on path from  $v$  to  $w$ :
      - longest edge is removed.
  - $O(\log n)$  time per edge.
- **Kruskal** as reference algorithm:
  1. Quicksort all  $m$  edges (offline);
  2. Insert edge iff its endpoints are in different components:
    - use union-find data structure.


# Online MSF on Random Graphs

- Random multigraphs:
  - $n = 4096$ , edges with random endpoints and weights;
  - expected diameter:  $O(\log n)$ ;
  - more edges  $\Rightarrow$  greater percentage of exposes.



# Online MSF on Augmented Stars


- Augmented stars:
  - $n = 65537$ , varying number (and length) of spokes;
  - diameter:  $O(65537/\text{\#spokes})$ ;
  - spoke edges, with length 1, processed first;
  - other edges (random), with length 2, processed later;
  - $10n$  edges in total.






# Online MSF on Augmented Stars

- Augmented stars:
  - $n = 65537$ , varying number (and length) of spokes;
  - diameter:  $O(65537/\text{\#spokes})$ .



# Online MSF and Cache Effects

- Procedure:
  - Partition vertices at random into  $n/32$  components of size 32;
  - Random edges: pick random component, then random pair within it.
  - Total number of edges:  $4n$ .




## Experiment: Shortest Paths

- Bellman's single source shortest path algorithm:
  - Assign **distance label** to each vertex.
    - \* initially,  $d(s) \leftarrow 0$  (source) and  $d(v) \leftarrow \infty$  (remaining vertices).
  - In each iteration, process all arcs  $(v, w)$  in fixed order:
    - \* if  $d(w) < d(v) + \ell(v, w)$ , **relax**  $(v, w)$ :
    - \* decrement  $d(w)$  by  $\Delta = d(v) + \ell(v, w) - d(w)$ .
  - Stop when an iteration does not relax any arc.
  - Running time:  $O(mn)$ .
- Dynamic trees:
  - Maintain the current shortest path tree and current  $d(\cdot)$ ;
  - When relaxing  $(v, w)$ , decrement  $d(\cdot)$  for all descendants of  $w$ ;
  - $O(mn \log n)$ , but may require fewer iterations.


## Experiment: Shortest Paths

- Random graph with Hamiltonian cycle:
  - $n$  arcs on Hamiltonian cycle with weight in  $[1;10]$ .
  - $3n$  random arcs with weight  $[1;1000]$ .
  - Edges processed in random order.



## Experiment: Shortest Paths

- Random graph with Hamiltonian cycle:
  - Dynamic trees reduce #iterations by 60% to 75%...
  - ...but increases running times by a factor of at least 50.



# Outline

- The Dynamic Trees problem
  - Existing data structures
  - A new worst-case data structure
  - A new amortized data structure
  - Experimental results
- ⇒ Final remarks

# Summary

- Main contributions:
  - new worst-case data structure;
  - new self-adjusting data structure:
    - \* uses contraction and path decomposition.
  - experimental analysis.
- Future work and open problems:
  - Worst-case data structure: do we really need Euler tours?
  - Self-adjusting data structure: can we make it worst-case?
  - Hybrid data structure?
  - Extend top trees?
  - Generalize top trees (grids, planar graphs, ...)?

**Thank You**



# Main Approaches

	Top Trees	ST-trees
principle	tree contraction	path decomposition
general interface	<b>YES</b>	<b>NO</b>
unrestricted trees	<b>YES</b>	<b>NO</b>
representation	level-based	recursive
overhead	<b>HIGH</b>	<b>LOW<sub>er</sub></b>